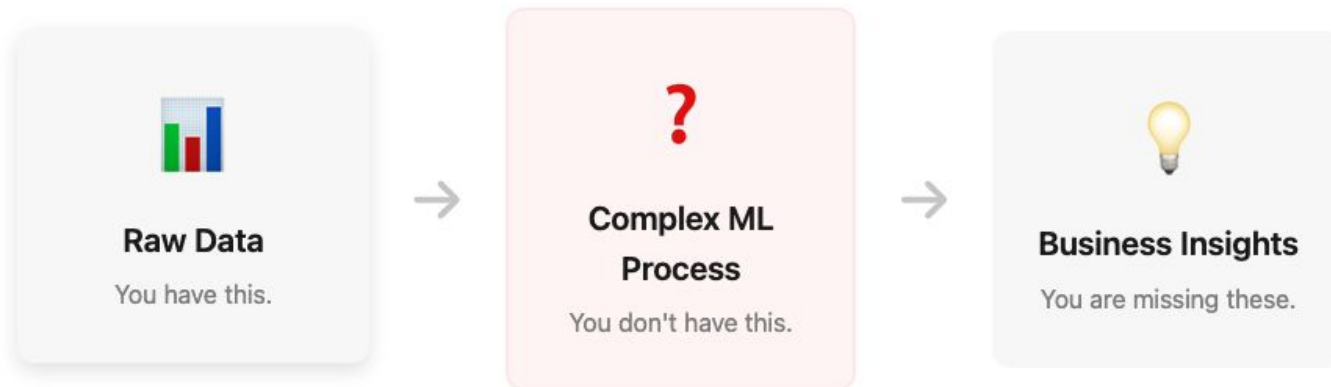


Universal ML Analysis Suite

From Data to Insights in Minutes

The Problem



Universal ML Analysis Suite

From Data to Insights in Minutes

Our Solution



CSV Upload

30 seconds



One-Click Analysis

5-10 minutes



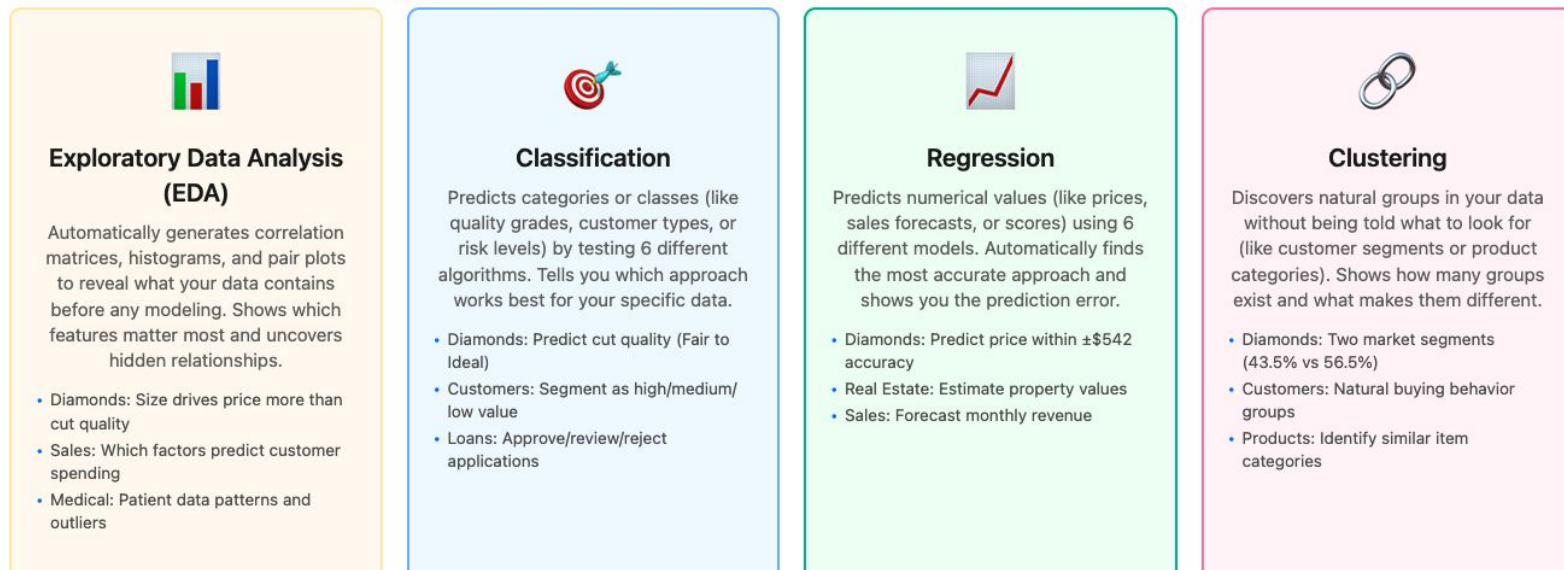
**Professional
Results**

Enterprise Quality

Two-Step Workflow



Suite Components



Diamond EDA Process



Diamond Dataset

53,940 samples

9 features: *carat, cut, color, clarity, depth, table, price, x, y, z*



Automated Analysis

5 minutes

*Correlation analysis
Distribution mapping
Pattern detection*




Key Discovery

Size = Price

*0.92 correlation reveals
carat weight drives pricing*

Correlation Matrix



Diamond Cut Classification

Predicting Quality Grades Automatically

Classification Challenge



Diamond Features

Physical Properties

*Carat, color, clarity, depth,
table, price, dimensions*



6 ML Algorithms

Automated Testing

*Logistic, KNN, Tree,
Forest,
SVC, Neural Network*



Cut Quality

5 Classes Predicted

*Fair, Good, Very Good,
Premium, Ideal*

Diamond Cut Classification

Predicting Quality Grades Automatically

Algorithm Performance Results

Neural Network

79.70%

Best overall accuracy

Random Forest

78.74%

Close second, stable

SVC

74.32%

Good performance

Decision Tree

71.43%

Fast training

KNN

53.75%

Moderate results

Logistic Regression

40.60%

Poor fit



CLASSIFICATION RESULTS SUMMARY

Algorithm	Accuracy	Training Time	Status
Logistic Regression	0.4060	0:00:00	✓ Success
Knn	0.5375	0:00:03	✓ Success
Decision Tree	0.7143	0:00:02	✓ Success
Random Forest	0.7874	0:01:16	✓ Success
Svc	0.7432	0:00:25	✓ Success
Ann	0.7970	0:01:13	✓ Success



Best Performer: Ann (0.7970 accuracy)



Classification completed successfully!

Diamond Price Prediction

98.15% Accuracy • \pm \$542 Average Error

Regression Challenge



Diamond Features

Physical & Quality

*Carat, cut, color, clarity,
depth, table, dimensions*



6 ML Models

Automated Testing

*Linear, KNN, Tree, Forest,
SVR, Neural Network*



Price Prediction

\$326 - \$18,823

*Mean: \$3,933
53,940 diamonds*

Diamond Price Prediction

98.15% Accuracy • ±\$542 Average Error

Model Performance Results

Random Forest

98.15% R^2

±\$542 RMSE

Exceptional accuracy

Neural Network

96.85% R^2

±\$708 RMSE

Strong performance

Decision Tree

96.66% R^2

±\$729 RMSE

Fast & interpretable

KNN

94.98% R^2

±\$894 RMSE

Good local patterns

Linear Regression

88.51% R^2

±\$1,351 RMSE

Simple baseline

SVR

51.45% R^2

±\$2,778 RMSE

Poor fit



REGRESSION RESULTS SUMMARY

Algorithm Status	R^2 Score	RMSE	Training Time
Linear Regression ✓ Success	0.8851	1351.26	0:00:00
Knn ✓ Success	0.9498	893.76	0:00:17
Decision Tree ✓ Success	0.9666	729.03	0:00:40
Random Forest ✓ Success	0.9815	542.02	0:00:16
Svr ✓ Success	0.5145	2778.03	0:00:48
Ann ✓ Success	0.9685	707.71	0:00:39



Best Performer: Random Forest ($R^2 = 0.9815$, RMSE = 542.02)



Regression completed successfully!

Diamond Market Segmentation

Discovering Natural Customer Groups

Clustering Challenge



Diamond Features

8 Dimensions

*Carat, color, clarity, depth,
table, dimensions (x, y, z)*



3 Algorithms

Pattern Discovery

*K-Means, Agglomerative,
Mean Shift clustering*



Market Segments

Natural Groups

*Two-tier market structure
43.5% vs 56.5% split*

Diamond Market Segmentation

Discovering Natural Customer Groups

Clustering Results & Market Structure

K-Means

0.3223

2 clusters • 100% coverage
Best silhouette score

Mean Shift

-0.0596

12 clusters • 27.8% coverage
Over-segmentation issue

Agglomerative

-0.1213

2 clusters • 18.5% coverage
Limited by sampling

Two-Tier Market Structure

Segment 1	23,485 (43.5%)
Segment 2	30,455 (56.5%)
Total Market	53,940 (100%)

CLUSTERING RESULTS SUMMARY

Algorithm	Clusters	Silhouette	Optimal
Param	Data	Coverage	Notes
Kmeans	2	0.3171	K=2
	100.0%	Full dataset	
Agglomerative	2	-0.1183	N=2
	18.5%	Sampled data	
Meanshift	12	-0.0596	BW=2.987
	27.8%	Sampled data	

★ Best performing algorithm based on silhouette score: Kmeans

Technology Stack

Tool Categories



Core Platform

Python + Scientific Stack

Python 3.12
pandas
numpy
scikit-learn



ML Algorithms

15 Models + 3 Categories

Logistic Regression
Linear Regression
K-Nearest Neighbors
Decision Tree
Random Forest
SVC
SVR
ANN



Visualization

5 Chart Types

matplotlib
seaborn
Correlation Matrix
Pair Plots
Histograms
Box Plots



Development

CLI + Web Interface

CLI interface
argparse
*Flask interface

Major Development Problems & Solutions



Repetitive Data Analysis

Same plots generated for every ML task. Output was 70% redundant EDA, 30% ML results.

Critical UX Issue



SOLVED



Two-Path Architecture

Created dedicated EDA script + clean ML outputs. Users choose exploration vs analysis.

Clean & Focused



Memory Limitations

53,940 samples crashed clustering algorithms. Agglomerative and Mean Shift became unusable.

Scalability Blocker



SOLVED



Smart Optimization

Intelligent sampling + garbage collection + progressive processing. Handles 53k+ samples efficiently.

Enterprise Scale

Marketing Strategies



SMB Market Focus

SMALL-MEDIUM BUSINESSES

Target the \$2B SMB market with "We have data but no data scientists" messaging. Position as enterprise analytics at startup prices.

LinkedIn Ads

Industry Webinars

Free Trials

Case Studies

Volume-based pricing



Consulting Partnerships

CONSULTING FIRMS

"Deliver client insights in hours, not weeks." Partner with analytics consultancies for white-label deployment and revenue sharing.

Partner Program

White Label

Revenue Share

Training

Custom pricing • Channel Growth



Vertical Expansion

INDUSTRY-SPECIFIC SOLUTIONS

Create templates for retail, healthcare, finance, real estate. "Diamond analysis" becomes "Customer analysis," "Property analysis," etc.

Industry Templates

Vertical Landing Pages

Trade Shows

Industry Partnerships

Custom Pricing • Premium Tiers